

Cervical Cancer Behavior Risk

February 26, 2023

1 Unsupervised Learning on Cervical Cancer Behavior Risk

The dataset was retrieved from the UC Irvine Machine Learning Repository. Courtesy goes to Dr. Sobar, Professor Rizanda Machmud and (Dr.) Adi Wijaya for providing the dataset to the repository.

Sobar, Machmud, R., & Wijaya, A. (2016). Cervical Cancer Behavior Risk Data Set. [Link](#)

1.1 Table of Contents

- [Background](#)
- Exploratory Data Analysis
- Model Development
 - [Summary](#)
- [Discussion](#)
- [Conclusion](#)

1.2 Background

Cancer is the uncontrollable growth of cells; cervical cancer is thus, the uncontrollable growth of cells in the cervix of women. According to the [CDC](#) page for cervical cancer, “human papillomavirus(HPV) is the main cause of cervical cancer. HPV is a common virus that is passed from one person to another during sex. At least half of sexually active people will have HPV at some point in their lives, but few women will get cervical cancer”.

Accordingly, there are screening tests to detect HPV and there is the HPV vaccine which can help with cervical cancer prevention. Screening tests include the pap smear (looks for precancers, cells changes) and the HPV test (looks for presence of HPV). When detected early on, the cancer is treatable.

The dataset provided by the UC Irvine Machine Learning Repository includes data from surveys conducted by the researchers, mentioned above, in Indonesia. The dataset is the result of the response of 72 respondents. The goal of the dataset was to attempt to predict cervical risk through behavioral measures. Sobar et al. used supervised learning methods, i.e. logistic regression, to determine cervical risk. More information can be found in their paper [here](#). As the title says, this notebook is going try to use unsupervised learning to see whether we can detect signs of cervical cancer risks through cluster models.

1.2.1 Goal

To reiterate, the goal of this notebook is going to be using unsupervised learning, more specifically clustering, to see whether we can differentiate groups that are susceptible to cervical cancer based on their behavioral determinants. The clustering model comparison will consist of 3 feature modifications: using the original data, using scaled data, and using post-PCA data. The models used will be hierarchical and K-means clustering.

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import AgglomerativeClustering, KMeans
from sklearn.metrics import confusion_matrix
```

1.3 Exploratory Data Analysis

```
[2]: data = pd.read_csv('sobar-72.csv')
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 72 entries, 0 to 71
```

```
Data columns (total 20 columns):
```

#	Column	Non-Null Count	Dtype
0	behavior_sexualRisk	72 non-null	int64
1	behavior_eating	72 non-null	int64
2	behavior_personalHygiene	72 non-null	int64
3	intention_aggregation	72 non-null	int64
4	intention_commitment	72 non-null	int64
5	attitude_consistency	72 non-null	int64
6	attitude_spontaneity	72 non-null	int64
7	norm_significantPerson	72 non-null	int64
8	norm_fulfillment	72 non-null	int64
9	perception_vulnerability	72 non-null	int64
10	perception_severity	72 non-null	int64
11	motivation_strength	72 non-null	int64
12	motivation_willingness	72 non-null	int64
13	socialSupport_emotionality	72 non-null	int64
14	socialSupport_appreciation	72 non-null	int64
15	socialSupport_instrumental	72 non-null	int64
16	empowerment_knowledge	72 non-null	int64
17	empowerment_abilities	72 non-null	int64
18	empowerment_desires	72 non-null	int64
19	ca_cervix	72 non-null	int64

```
dtypes: int64(20)
memory usage: 11.4 KB
```

Initial inspection shows 72 data points and none of the columns seem to have null values. There are 20 columns total. The goal is going to be trying to use unsupervised learning to identify the ca_cervix column, i.e. the presence or absence of cervical cancer.

According to the background information provided by the dataset, there are 19 attributes that come from 8 variables. The variables is written as the first part of the attribute name - behavior, intention, attitude, norm, perception, motivation, socialSupport, and empowerment.

```
[3]: data.describe()
```

```
[3]:
```

	behavior_sexualRisk	behavior_eating	behavior_personalHygiene	\
count	72.000000	72.000000	72.000000	
mean	9.666667	12.791667	11.083333	
std	1.186782	2.361293	3.033847	
min	2.000000	3.000000	3.000000	
25%	10.000000	11.000000	9.000000	
50%	10.000000	13.000000	11.000000	
75%	10.000000	15.000000	14.000000	
max	10.000000	15.000000	15.000000	

	intention_aggregation	intention_commitment	attitude_consistency	\
count	72.000000	72.000000	72.000000	
mean	7.902778	13.347222	7.180556	
std	2.738148	2.374511	1.522844	
min	2.000000	6.000000	2.000000	
25%	6.000000	11.000000	6.000000	
50%	10.000000	15.000000	7.000000	
75%	10.000000	15.000000	8.000000	
max	10.000000	15.000000	10.000000	

	attitude_spontaneity	norm_significantPerson	norm_fulfillment	\
count	72.000000	72.000000	72.000000	
mean	8.611111	3.125000	8.486111	
std	1.515698	1.845722	4.907577	
min	4.000000	1.000000	3.000000	
25%	8.000000	1.000000	3.000000	
50%	9.000000	3.000000	7.000000	
75%	10.000000	5.000000	14.000000	
max	10.000000	5.000000	15.000000	

	perception_vulnerability	perception_severity	motivation_strength	\
count	72.000000	72.000000	72.000000	
mean	8.513889	5.388889	12.652778	
std	4.275686	3.400727	3.207209	
min	3.000000	2.000000	3.000000	

25%	5.000000	2.000000	11.000000
50%	8.000000	4.000000	14.000000
75%	13.000000	9.000000	15.000000
max	15.000000	10.000000	15.000000

	motivation_willingness	socialSupport_emotionality \
count	72.000000	72.000000
mean	9.694444	8.097222
std	4.130406	4.243171
min	3.000000	3.000000
25%	7.000000	3.000000
50%	11.000000	9.000000
75%	13.000000	11.250000
max	15.000000	15.000000

	socialSupport_appreciation	socialSupport_instrumental \
count	72.000000	72.000000
mean	6.166667	10.375000
std	2.897303	4.316485
min	2.000000	3.000000
25%	3.750000	6.750000
50%	6.500000	12.000000
75%	9.000000	14.250000
max	10.000000	15.000000

	empowerment_knowledge	empowerment_abilities	empowerment_desires \
count	72.000000	72.000000	72.000000
mean	10.541667	9.319444	10.277778
std	4.366768	4.181874	4.482273
min	3.000000	3.000000	3.000000
25%	7.000000	5.000000	6.750000
50%	12.000000	10.000000	11.000000
75%	15.000000	13.000000	15.000000
max	15.000000	15.000000	15.000000

	ca_cervix
count	72.000000
mean	0.291667
std	0.457719
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

```
[4]: data.head()
```

```

[4]:  behavior_sexualRisk  behavior_eating  behavior_personalHygiene  \
0          10          13          12
1          10          11          11
2          10          15           3
3          10          11          10
4           8          11           7

      intention_aggregation  intention_commitment  attitude_consistency  \
0              4              7              9
1             10             14              7
2              2             14              8
3             10             15              7
4              8             10              7

      attitude_spontaneity  norm_significantPerson  norm_fulfillment  \
0              10              1              8
1              7              5              5
2             10              1              4
3              7              1              5
4              8              1              5

      perception_vulnerability  perception_severity  motivation_strength  \
0              7              3             14
1              4              2             15
2              7              2              7
3              4              2             15
4              3              2             15

      motivation_willingness  socialSupport_emotionality  \
0              8              5
1             13              7
2              3              3
3             13              7
4              5              3

      socialSupport_appreciation  socialSupport_instrumental  \
0              7             12
1              6              5
2              6             11
3              4              4
4              6             12

      empowerment_knowledge  empowerment_abilities  empowerment_desires  \
0              12             11              8
1              5              4              4
2              3              3             15
3              4              4              4

```

4

5

4

7

	ca_cervix
0	1
1	1
2	1
3	1
4	1

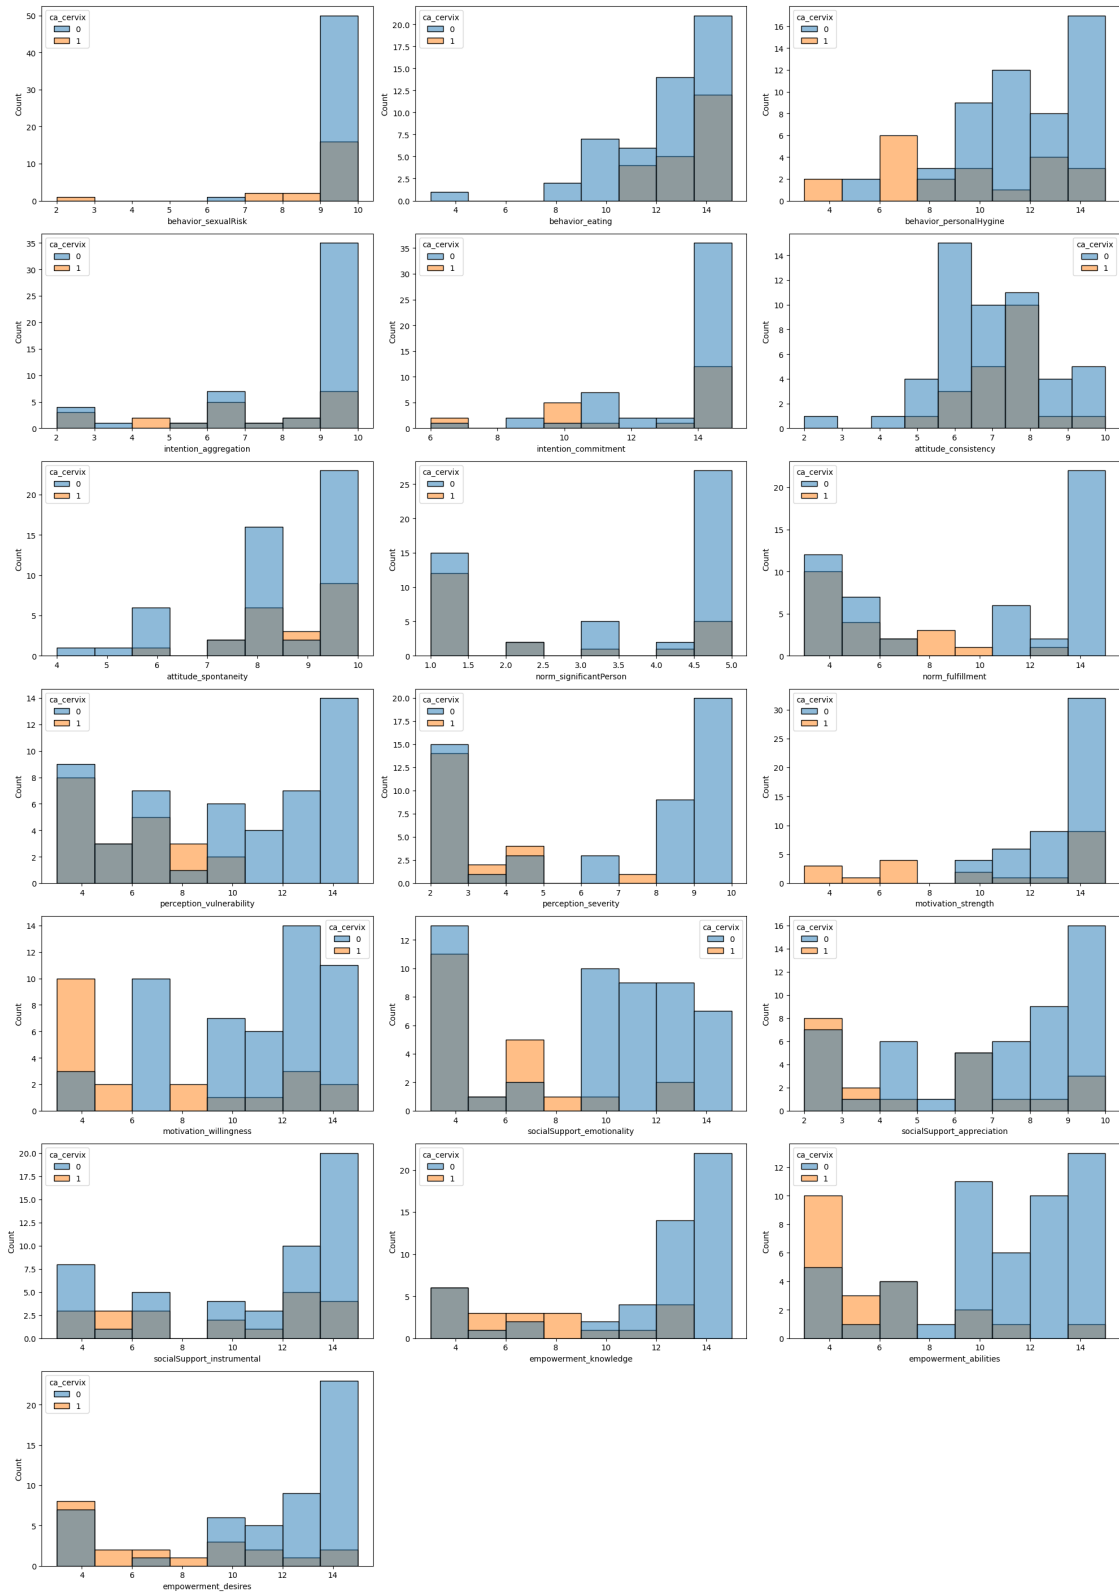
Looking at the description and a subset of the dataset, the values are related to each respondents' answers to the questionnaire. But, it isn't clear what the values represent in terms of context. It is unclear what the measure of the values are; are they representing awareness, their thoughts towards each attribute, or something else?

```
[5]: # removing the classification col for attribute eda.
columns = list(data.columns.values)
columns.remove('ca_cervix')

fig = plt.figure(figsize=(20, 40))
rows, cols = 10, 3

# subplots for each attribute
for i, col in enumerate(columns):
    ax = fig.add_subplot(rows, cols, i+1)
    sns.histplot(data=data, x=col, hue='ca_cervix')
    ax.set_xlabel(col)

fig.tight_layout()
plt.show()
```



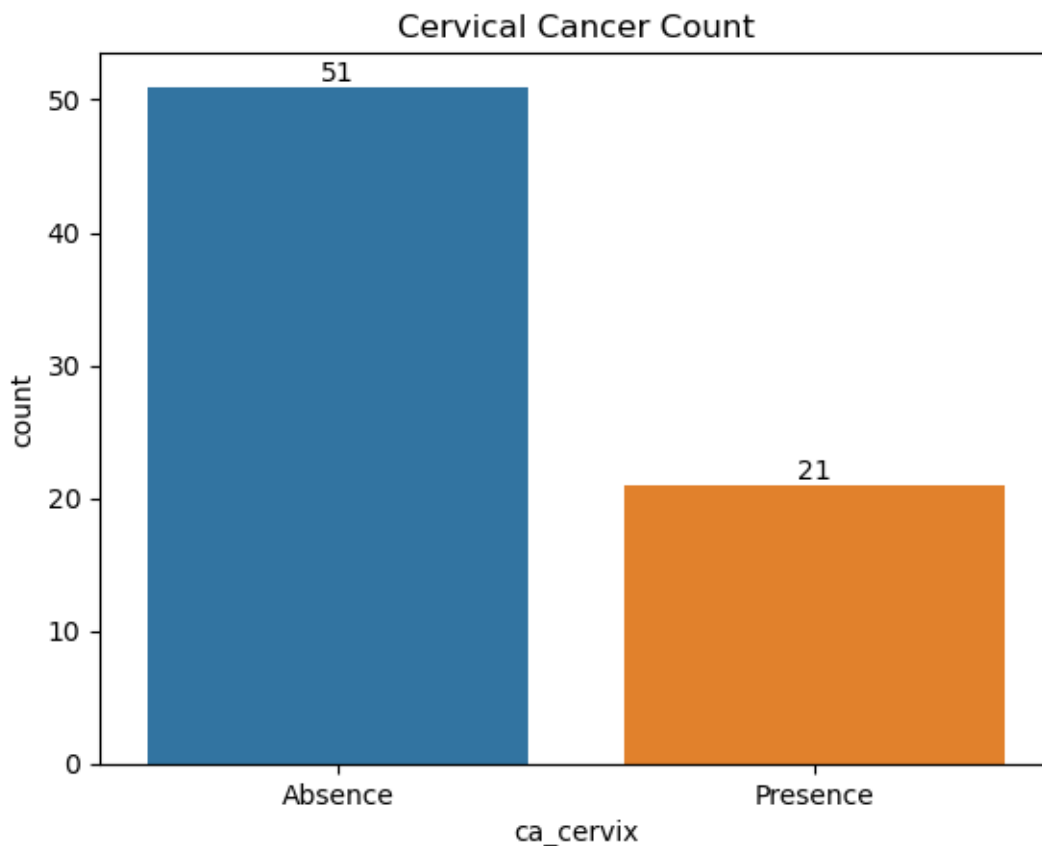
In the plots above, we can see that there is a range of data and the attributes don't look normal. I think it's necessary for it to be normal for unsupervised learning methods. From the histogram, we also notice some points that could be outliers; for example in the `behavior_sexualRisk` plot, most points are under 10 whereas there's a lone point at 2. The same pattern can be seen in the `intention_commitment` plot. We are not going to remove the outliers however, because they may be useful in the clustering differentiation process.

We do see some small trends for example in the `perception_vulnerability` chart, we see that respondents with cervical cancer tend to have mid-low answers to the related questions and none of the higher answers has cervical cancer. There is also the `empowerment_abilities` plot that shows low scores tend to have the cancer.

```
[6]: ax = sns.countplot(data=data, x='ca_cervix')


# find the count labels
for i in ax.containers:
    ax.bar_label(i,)

ax.set_xticklabels(['Absence', 'Presence'])
plt.title('Cervical Cancer Count')
plt.show()
```




```
[7]: corr = data.corr()
# used to take the top half of the correlation chart away for easier viewing
mask = np.triu(np.ones_like(corr, dtype=bool))

plt.figure(figsize=(20,20))
sns.heatmap(corr, linewidths=0.05, center=0, cmap='coolwarm', mask=mask,
            square=True, annot=True, fmt=".2f")
plt.title('Correlation Chart')
plt.show()
```




We see from the correlation plot that some attributes are highly correlated with each other. For example, perception_severity and norm_fulfillment are highly correlated at 0.85. Another notable one is empowerment_abilities and empowerment_knowledge at 0.84. This one you can reason out that the more knowledge someone has on the cancer, the more they feel they are in control of

preventing it.

In the plot below is the correlation of cervical cancer against all the attributes. We see that the highest correlation is a negative correlation with empowerment_abilities at -0.54.

```
[8]: # correlation with the presence/absence of cancer
data.corr()['ca_cervix'][: -1].sort_values().plot(kind='bar')
plt.title('Correlation with Cervix Cancer')
plt.show()
```



Because we are going to be using unsupervised learning methods, the apparent relationships we see from the plots above is good to note but won't be useful in the clustering process. Unsupervised learning is trying to find latent patterns to separate the respondents into groups - presence of

absence of cervical cancer.

Because the labels are given to us, we are going to use it to compare with the cluster models.

Below we scale the data between 0 and 1 using the standard scaler function. This will be the scaled features used for modeling. I choose to scale the data because as seen in the description of the dataset, the features have different min and max values. This corresponds to questions that had answers ranging from 0 to 10, 0 to 15, etc. Scaling the features would put them all on the same footing for model development.

```
[9]: features = data.drop(['ca_cervix'], axis=1)
```

```
[10]: scaler = StandardScaler()
      scaler.fit(features)


      features_scaled = pd.DataFrame(scaler.transform(features), columns=features.
      ↪columns)
```

Finally, all that's left is the PCA feature set. Why choose to do a PCA when the number of features isn't overly large or sparse? I chose to perform a PCA feature just to reduce the number of dimensions and also to try to reduce the correlation present between the features which may lead to some unknown interactions. Instead PCA will just create a number of features that explains the most variance in the dataset.

```
[11]: explained_var = []
      for i in range(1,20):
          pca = PCA(n_components=i).fit(features_scaled)
          explained_var.append(sum(pca.explained_variance_ratio_))
```

The chosen cutoff for PCA is going to be around 90% explained variance. This seems to be around 10-features so that's what is going to be used for the PCA features part of the model development.

```
[12]: ax = sns.scatterplot(x=range(1, 20), y=explained_var)
      ax.set_xticks(range(len(explained_var)), labels=range(0,19))
      ax.axhline(.9, color='red', dashes=(5, 2))
      plt.xlabel('n-components')
      plt.ylabel('Explained Variance')
      plt.show()
```



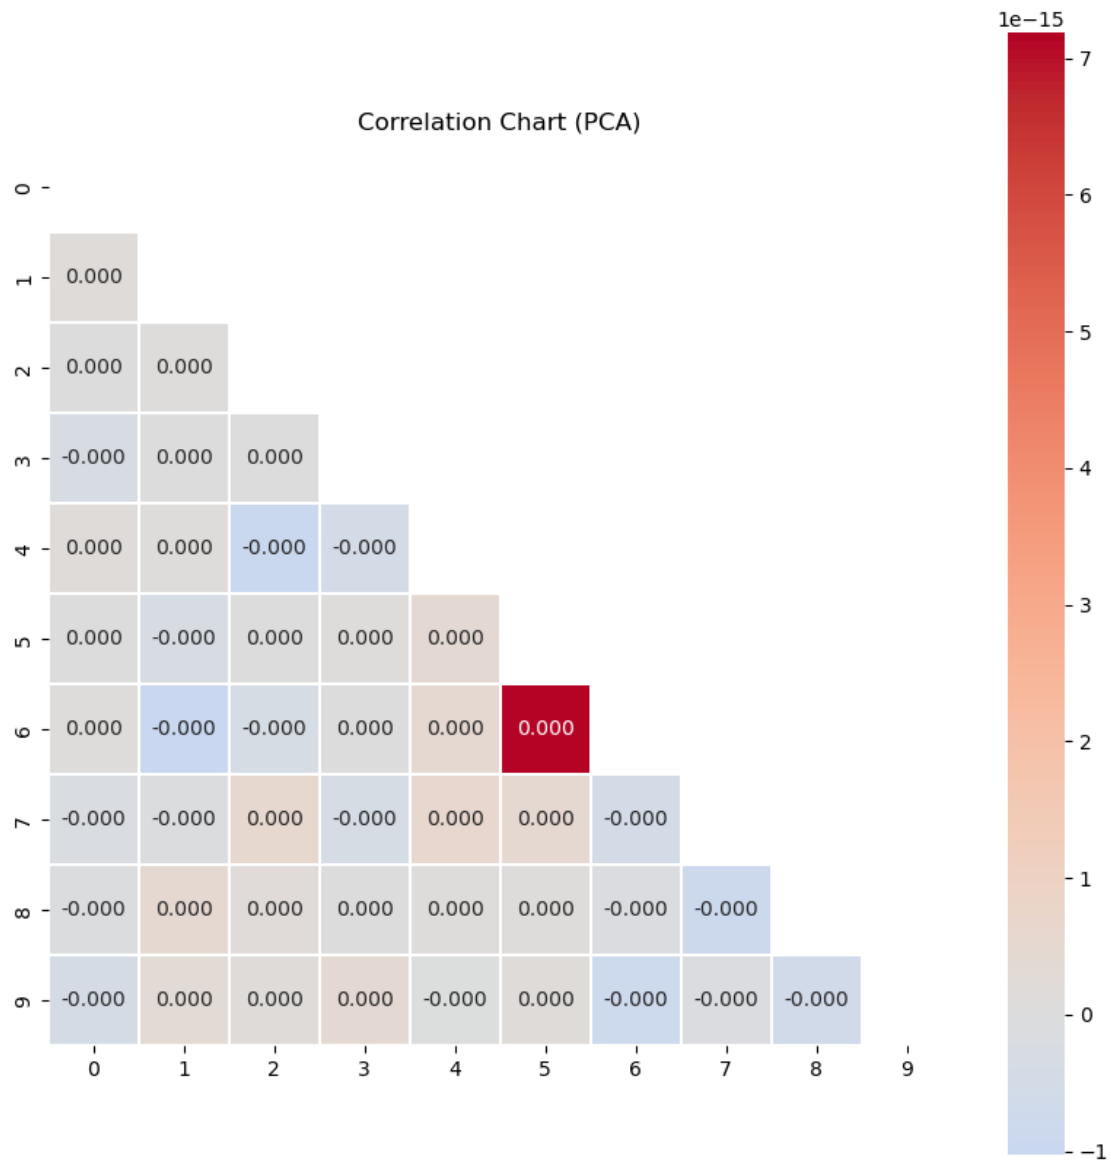
```
[13]: pca = PCA(n_components=10).fit(features_scaled)
pca_features = pca.transform(features_scaled)
```

```
[14]: pca_features[0]
```

```
[14]: array([ 0.34259186, -0.93663453,  0.44628193,  1.25042552,  2.13653477,
            1.04437742, -1.49113714, -1.74232282, -0.24873946,  0.98933645])
```

```
[15]: pca_df = pd.DataFrame(pca_features)
```

```
[16]: corr_pca = pca_df.corr()
mask = np.triu(np.ones_like(corr_pca, dtype=bool))
plt.figure(figsize=(10,10))
sns.heatmap(corr_pca, linewidths=0.05, center=0, cmap='coolwarm', mask=mask,
            square=True, annot=True, fmt=".3f")
plt.title('Correlation Chart (PCA)')
plt.show()
```



You see that after the PCA, the features chosen have basically 0 correlation with each other.

1.4 Model Development

1.4.1 Using Original Data

Running an initial hierarchical clustering method. 'n_clusters' is set at 2 because that's what we are hoping to separate into presence and absence of cancer.

```
[17]: # intial test run of the model
model = AgglomerativeClustering(n_clusters=2, linkage='average').fit(features)
model.labels_
```


```
[17]: array([0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0,
          0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0,
          0, 0, 0, 0, 0, 0], dtype=int64)
```

```
[18]: n = len(data)

# compares the model labels against the dataset classification
print('accuracy:', sum(model.labels_ == np.array(data.ca_cervix)) / n)

sns.heatmap(confusion_matrix(y_pred=model.labels_, y_true=np.array(data.
    ↪ca_cervix)), annot=True)
plt.show()
```

accuracy: 0.8194444444444444



We have here that the accuracy of the clusters to the labels are around 82%. This is pretty decent. Looking at the confusion matrix we see that the cluster does relatively well. It has a precision of $(16/24) = 67\%$ and a true positive rate of $(16/21) = 76\%$ and a false negative $(5/21) = 24\%$.


Cancer in general is a detrimental disease to have, thus it's important to discern whether a person has it or not. As a result, it's important to look at the accuracy of the model and the false negative rate and hopefully have the FNR as low as possible.

Next, we are going to see whether an adjustment to the hierarchical cluster model can improve those metrics.


```
[19]: # tests given iteration of agglomerative clustering and prints combination/  
      ↪ accuracy and confusion matrix.  
      # prints highest and increasing accuracy only.  
      def agcluster(feature_df, n):  
          acc=0  
          for i in ['ward', 'complete', 'average', 'single']:  
              for j in ['euclidean', 'manhattan', 'l1', 'l2', 'cosine']:  
                  try:  
                      mod = AgglomerativeClustering(n_clusters=2, linkage=i, ↪  
          ↪ affinity=j).fit(feature_df)  
                      total = max(sum(mod.labels_ == np.array(data.ca_cervix)), ↪  
          ↪ sum(mod.labels_ != np.array(data.ca_cervix)))  
                      acc_mod = total/n  
                      if acc_mod >= acc:  
                          acc = acc_mod  
                          print(i, j, round(acc, 5))  
                          sns.heatmap(confusion_matrix(y_pred=mod.labels_, y_true=np.  
          ↪ array(data.ca_cervix)), annot=True)  
                          plt.ylabel('True Labels')  
                          plt.xlabel('Predicted Labels')  
                          plt.show()  
  
                  except:  
                      continue
```

```
[20]: agcluster(features, n)
```


```
ward euclidean 0.83333
```



average manhattan 0.83333



average 11 0.83333



We see that the linkage-metric combination that yielded the highest accuracy is [ward, euclidean], [average, manhattan], and [average, l1] at 83%. We see that these models yield the same result with 1 less false negative (19% FNR) compared to the initial model.


Next, we are going to try the K-means clustering method to see whether there is any differences.

```
[21]: # initial kmeans test run.
kmeans = KMeans(n_clusters=2, random_state=41).fit(features)

# max was done here to find whether the 0 label = 0 or 0 label = 1.
print('accuracy:', max(sum(kmeans.labels_ == data.ca_cervix), sum(kmeans.
    ↳ labels_ != data.ca_cervix)) / n)

sns.heatmap(confusion_matrix(y_pred=kmeans.labels_, y_true=np.array(data.
    ↳ ca_cervix)), annot=True)
plt.show()
```

accuracy: 0.7777777777777778



- Accuracy: 78%
- FNR: 19%

The initial K-means model yielded a lower accuracy, 5% less, than the best hierarchical model.


Let's see if there can be any further adjustments to be made.

```
[22]: # prints kmeans combination/accuracy and confusion matrix
# prints only the best and increasingly higher accuracy models.
def kmeanscluster(features_df, n):
    acc = 0
    for i in ['auto', 'elkan']:
        for j in ['k-means++', 'random']:
            km_mod = KMeans(n_clusters=2, init=j, algorithm=i, random_state=41).
            fit(features_df)
            total = max(sum(km_mod.labels_ == np.array(data.ca_cervix)),
            sum(km_mod.labels_ != np.array(data.ca_cervix)))
            acc_mod = total/n
            if acc_mod >= acc:
                acc = acc_mod
                print(i, j, round(acc, 5))
                sns.heatmap(confusion_matrix(y_pred=km_mod.labels_, y_true=np.
            array(data.ca_cervix)), annot=True)
```


```
plt.ylabel('True Labels')
plt.xlabel('Predicted Labels')
plt.show()
```

```
[23]: kmeanscluster(features, n)
```


auto k-means++ 0.77778




auto random 0.77778



elkan k-means++ 0.77778



elkan random 0.77778




Looks like nothing changed even with different parameters.


1.4.2 Using Scaled Features

```
[24]: agcluster(features_scaled, n)
```


```
ward euclidean 0.83333
```



complete manhattan 0.83333




complete 11 0.83333




Nothing changed from the hierarchical model.

```
[25]: kmeanscluster(features_scaled, n)
```


```
auto k-means++ 0.79167
```




auto random 0.79167



elkan k-means++ 0.79167



elkan random 0.79167

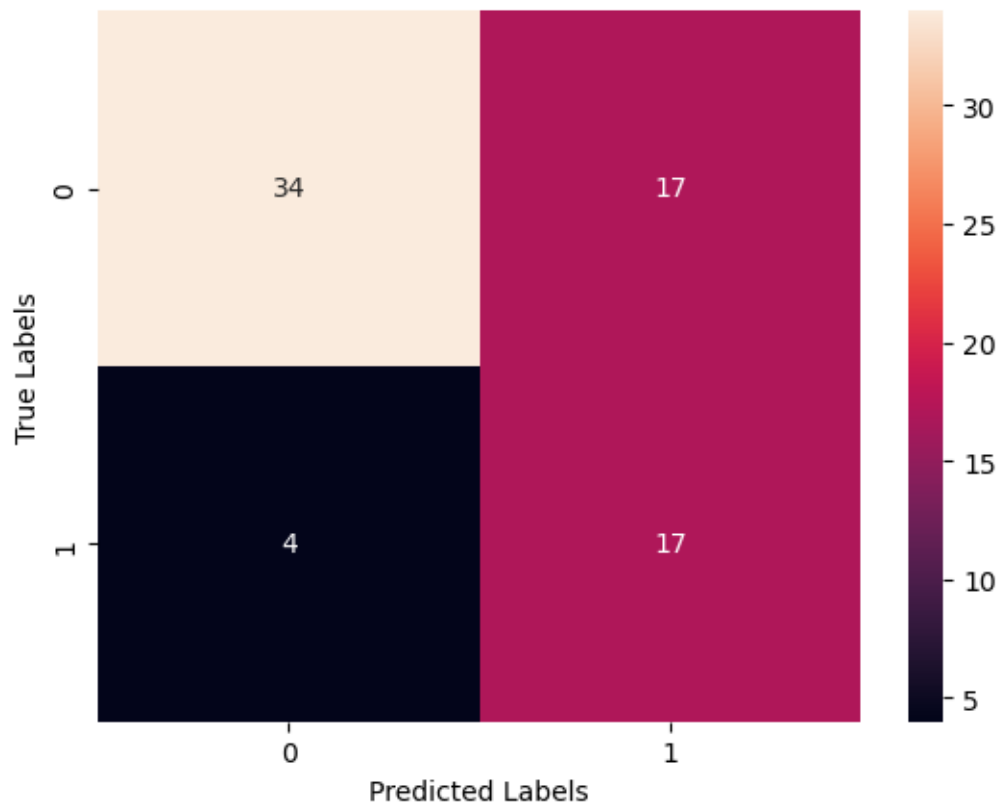


The K-means model does see a very small improvement in their accuracy - 78% to 79%. The FNR remains the same.


1.4.3 Using PCA Features

```
[26]: agcluster(pca_features, n)
```


```
ward euclidean 0.70833
```




complete euclidean 0.72222




complete manhattan 0.72222




complete 11 0.72222



complete 12 0.72222




complete cosine 0.79167

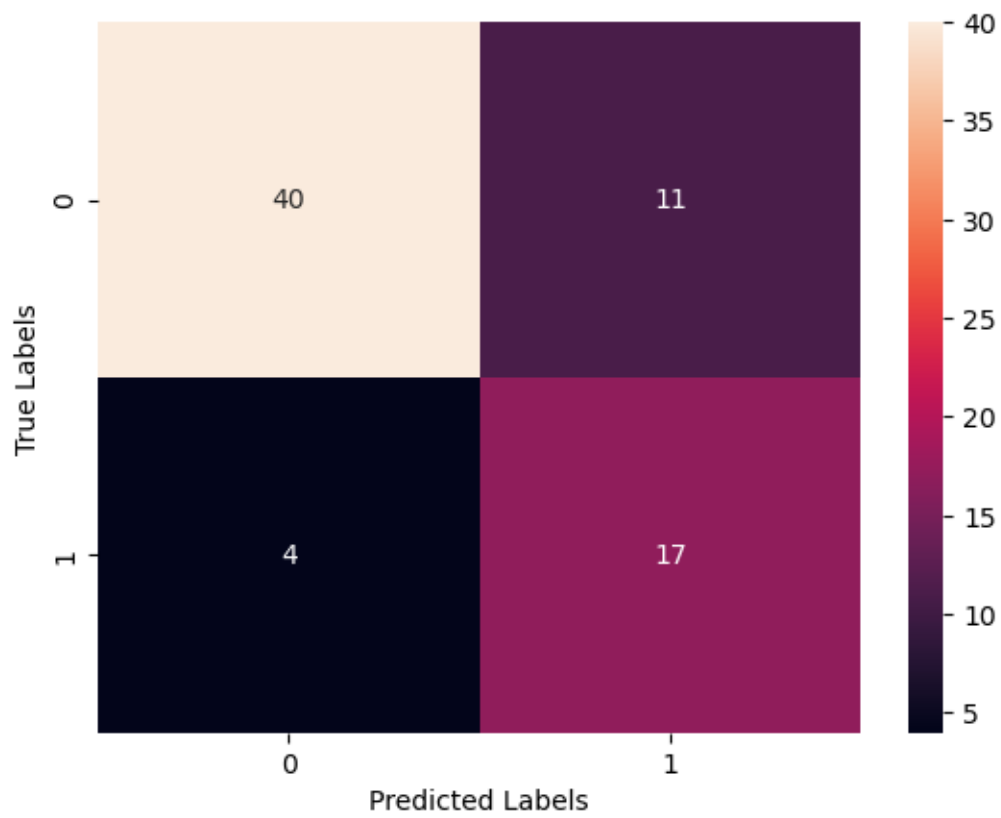


```
[27]: kmeanscluster(pca_features, n)
```


```
auto k-means++ 0.79167
```



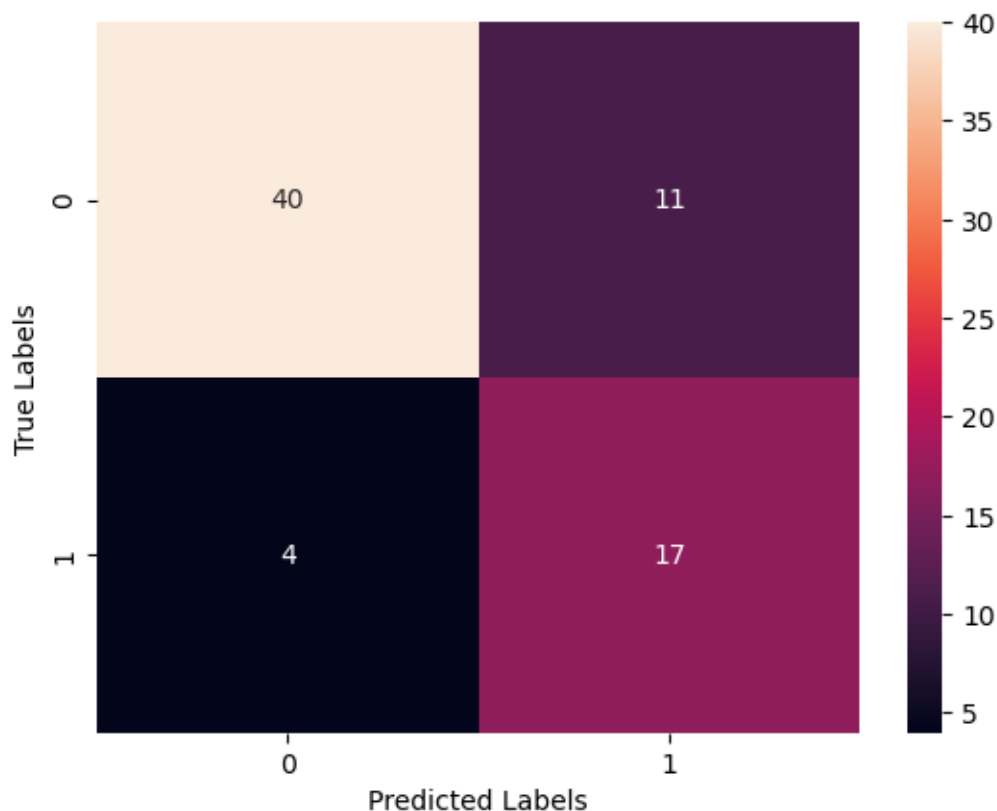
auto random 0.79167



elkan k-means++ 0.79167



elkan random 0.79167



1.4.4 Summary

Method	Original Features	Scaled Features	PCA Features
Hierarchical Clustering	0.833	0.833	0.792
KMeans Clustering	0.778	0.792	0.792

In summary we see that scaling the features increased the K-means result and that performing pca on the scaled features yielded the same result. In contrast, the hierarchical clustering yielded the highest accuracy on the original dataset and after scaling it. The FNR is not shown but it was the same for all best models under each category at 19%.

The hierarchical clustering yielded the highest accuracy using the original dataset probably because there was more context involved in the clustering which helped group certain points together. K-means didn't do as well there because the scaling was different very some features which may have caused some points to be further away then should be. Also K-means is highly susceptible to the effects of outliers and remember, there were some points that were outliers which may have caused the mean centroids to shift in certain directions that would miss some points. You do see it improve after scaling.

In general, it would be best to scale the features for clustering and in this case, hierarchical clustering is the best model.

1.5 Discussion

Cervical cancer is one of the cancers that when found early, can be treatable. It is important to find methods to be able to detect it early on. Previously mentioned, there are screening tests that can help detect the presence of the HPV itself and signs of cell changing but these are pretty much after the fact. If we were able to use behavior determinants to find people at risk of cervical cancer, preventative measures can be given to the people.

As we see from the results, clustering can help determine the presence/absence of cervical cancer with an overall accuracy of 83% and a false negative rate of 19%. It is a decent model and can be used as a possible precursory test to find people at risk of cervical cancer.

It is important to consider that the data was obtained via a questionnaire. Questionnaires are subjective and the phrasing of questions may cause different understanding of what is asked of the respondents. We also don't know what was asked for each behavioral determinant in the questionnaire. Another consideration is the location of the survey. The dataset values originally came from responses from Indonesian women. Not to say they are not important, but due to the culture, the data may not be representative for the entire women population.

For future works, hopefully there would be more data available and perhaps more variety in the data. Variety in terms of the nationality of the women so we can see a more representative response and thus a more representative model.

1.6 Conclusion

Hierarchical modeling is a viable method for determining women at risk of cervical cancer based on their behavioral determinants. The model yielded a decent accuracy of 83% and a false negative rate of 19%. It seems relatively high for some but the result came from a dataset with 72 points and fewer examples of true positives. There are some improvements that could be done but in the context of women in Indonesia, this hierarchical modeling is a potential screening method for cervical cancer.